

Assignment no 2 :

File Management tasks in Hadoop

1. Create a directory in HDFS at given path(s).

Usage:

```
hadoop fs -mkdir <paths>
```

Example:

```
hadoop fs -mkdir /user/saurzcode/dir1 /user/saurzcode/dir2
```

2. List the contents of a directory.

Usage :

```
hadoop fs -ls <args>
```

Example:

```
hadoop fs -ls /user/saurzcode
```

3. Upload and download a file in HDFS.

Upload:

hadoop fs -put:

Copy single src file, or multiple src files from local file system to the Hadoop data file system

Usage:

```
hadoop fs -put <localsrc> ... <HDFS_dest_Path>
```

Example:

```
hadoop fs -put /home/saurzcode/Samplefile.txt /user/saurzcode/dir3/
```

Download: hadoop fs -get:

Copies/Downloads files to the local file system

Usage:

```
hadoop fs -get <hdfs_src> <localdst>
```

Example:

```
hadoop fs -get /user/saurzcode/dir3/Samplefile.txt /home/
```

4. See contents of a file

Same as unix cat command:

Usage:hadoop fs -cat <path[filename]>

Example:

```
hadoop fs -cat /user/saurzcode/dir1/abc.txt
```

5. Copy a file from source to destination

This command allows multiple sources as well in which case the destination must be a directory.

Usage:

```
hadoop fs -cp <source> <dest>
```

Example:

```
hadoop fs -cp /user/saurzcode/dir1/abc.txt /user/saurzcode/dir2
```

6. Copy a file from/To Local file system to HDFS

copyFromLocal

Usage:

```
hadoop fs -copyFromLocal <localsrc> URI
```

Example:

```
hadoop fs -copyFromLocal /home/saurzcode/abc.txt /user/saurzcode/abc.txt
```

Similar to put command, except that the source is restricted to a local file reference.

copyToLocal

Usage:

```
hadoop fs -copyToLocal [-ignorecrc] [-crc] URI <localdst>
```

Similar to get command, except that the destination is restricted to a local file reference.

7. Move file from source

to destination.

Note:- Moving files across filesystem is not permitted.

Usage :

```
hadoop fs -mv <src> <dest>
```

Example:

```
hadoop fs -mv /user/saurzcode/dir1/abc.txt /user/saurzcode/dir2
```

8. Remove a file or directory in HDFS.

Remove files specified as argument. Deletes directory only when it is empty

Usage :

```
hadoop fs -rm <arg>
```

Example:

```
hadoop fs -rm /user/saurzcode/dir1/abc.txt
```

Recursive version of delete.

Usage :

```
hadoop fs -rmr <arg>
```

Example:

```
hadoop fs -rmr /user/saurzcode/
```

9. Display last few lines of a file. Similar to tail command in Unix.

Usage :

```
hadoop fs -tail <path[filename]>
```

Example:

```
hadoop fs -tail /user/saurzcode/dir1/abc.txt
```

10. Display the aggregate length of a file.

Usage :

```
hadoop fs -du <path>
```

Example:

```
hadoop fs -du /user/saurzcode/dir1/abc.txt
```

Assignment 3

Word Count Map Reduce program to understand Map Reduce Paradigm

Source code:

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.fs.Path;
public class WordCount
{
    public static class Map extends Mapper<LongWritable,Text,Text,IntWritable> {
        public void map(LongWritable key, Text value,Context context) throws
        IOException,InterruptedException{
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                value.set(tokenizer.nextToken());
                context.write(value, new IntWritable(1));
            }
        }
    }

    public static class Reduce extends Reducer<Text,IntWritable,Text,IntWritable> {
        public void reduce(Text key, Iterable<IntWritable> values,Context context)
        throws IOException,InterruptedException {
            int sum=0;
            for(IntWritable x: values)
            {
                sum+=x.get();
            }
            context.write(key, new IntWritable(sum));
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf= new Configuration();
        Job job = new Job(conf, "My Word Count Program");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
```

```

Path outputPath = new Path(args[1]);
//Configuring the input/output path from the filesystem into the job
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
//deleting the output path automatically from hdfs so that we don't have to
delete it explicitly
outputPath.getFileSystem(conf).delete(outputPath);
//exiting the job only if the flag value becomes false
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

The entire MapReduce program can be fundamentally divided into three parts:

- Mapper Phase Code
- Reducer Phase Code
- Driver Code

We will understand the code for each of these three parts sequentially.

Mapper code:

```

public static class Map extends
Mapper<LongWritable,Text,Text,IntWritable> {
public void map(LongWritable key, Text value, Context context)
throws IOException,InterruptedException {
String line = value.toString();
StringTokenizer tokenizer = new StringTokenizer(line);
while (tokenizer.hasMoreTokens()) {
value.set(tokenizer.nextToken());
context.write(value, new IntWritable(1));
}
}
}

```

- We have created a class Map that extends the class Mapper which is already defined in the MapReduce Framework.
- We define the data types of input and output key/value pair after the class declaration using angle brackets.
- Both the input and output of the Mapper is a key/value pair.
- Input:
 - The key is nothing but the offset of each line in the text file: LongWritable
 - The value is each individual line (as shown in the figure at the right): Text
- Output:
 - The key is the tokenized words: Text
 - We have the hardcoded value in our case which is

1: IntWritable

- Example – Dear 1, Bear 1, etc.
- We have written a java code where we have tokenized each word and assigned them a hardcoded value equal to 1.

Reducer Code:

```
public static class Reduce extends
Reducer<Text,IntWritable,Text,IntWritable> {
public void reduce(Text key, Iterable<IntWritable>
values,Context context)
throws IOException,InterruptedException {
int sum=0;
for(IntWritable x: values)
{
sum+=x.get();
}
context.write(key, new IntWritable(sum));
}}
```

- We have created a class Reduce which extends class Reducer like that of Mapper.
 - We define the data types of input and output key/value pair after the class declaration using angle brackets as done for Mapper.
 - Both the input and the output of the Reducer is a keyvalue pair.
- Input:

- The key nothing but those unique words which have been generated after the sorting and shuffling phase: Text
- The value is a list of integers corresponding to each key:

IntWritable

- Example – Bear, [1, 1], etc.
- Output:
 - The key is all the unique words present in the input text file: Text
 - The value is the number of occurrences of each of the unique words: IntWritable
- Example – Bear, 2; Car, 3, etc.
- We have aggregated the values present in each of the list corresponding to each key and produced the final answer.
- In general, a single reducer is created for each of the unique words, but, you can specify the number of reducer in mapred site.xml.

Driver Code:

```

Configuration conf= new Configuration();
Job job = new Job(conf,"My Word Count Program");
job.setJarByClass(WordCount.class);
job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
Path outputPath = new Path(args[1]);
//Configuring the input/output path from the filesystem into the
job FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

```

- In the driver class, we set the configuration of our MapReduce job to run in Hadoop.
- We specify the name of the job , the data type of input/ output of the mapper and reducer.
- We also specify the names of the mapper and reducer classes. • The path of the input and output folder is also specified. • The method setInputFormatClass () is used for specifying that how a Mapper will read the input data or what will be the unit of work. Here, we have chosen TextInputFormat so that single line is read by the mapper at a time from the input text file.
- The main () method is the entry point for the driver. In this method, we instantiate a new Configuration object for the job.

Run the MapReduce code:

The command for running a MapReduce code is:

```

hadoop jar hadoop-mapreduce-example.jar WordCount /
sample/input /sample/output

```

Assignment 4.

Weather Report POC-Map Reduce Program to analyse time-temperature statistics and

generate report with max/min temperature.

Problem Statement:

1. The system receives temperatures of various cities(Austin, Boston,etc) of USA captured at regular intervals of time on each day in an input file. 2. System will process the input data file and generates a report with Maximum and Minimum temperatures of each day along with time. 3. Generates a separate output report for each city.

Ex: Austin-r-00000

Boston-r-00000

Newjersy-r-00000

Baltimore-r-00000

California-r-00000

Newyork-r-00000

Expected output:- In each output file record should be like this:

25-Jan-2014 Time: 12:34:542 MinTemp: -22.3 Time: 05:12:345 MaxTemp: 35.7

First **download input file** which contains temperature statistics with time for multiple cities. Schema of record set : CA_25-Jan-2014 00:12:345 15.7 01:19:345 23.1 02:34:542 12.3

CA is city code, here it stands for California followed by date. After that each pair of values represent time and temperature.

Mapper class and map method: -

The very first thing which is required for any map reduce problem is to understand what will be the type of keyIn, ValueIn, KeyOut, ValueOut for the given Mapper class and followed by type of map method parameters. •

public class WhetherForecastMapper extends Mapper <Object, Text, Text, Text>

- Object (keyIn) - Offset for each line, line number 1, 2...
- Text (ValueIn) - Whole string for each line (CA_25-Jan-2014 00:12:345)
- Text (KeyOut) - City information with date information as string • Text (ValueOut) - Temperature and time information which need to be passed to reducer as string.
- public void map(**Object keyOffset, Text dayReport, Context con**) { }
- *KeyOffset* is like line number for each line in input file.
- *dayreport* is input to map method - whole string present in one line of input file.

- *con* is context where we write mapper output and it is used by reducer.

Reducer class and reducer method:-

Similarly, we have to decide what will be the type of keyIn, ValueIn, KeyOut, ValueOut for the given Reducer class and followed by type of reducer method parameters.

- public class WhetherForcastReducer extends Reducer<**Text, Text, Text, Text**>
- Text(keyIn) - it is same as keyOut of Mapper.
- Text(ValueIn)- it is same as valueOut of Mapper.
- Text(KeyOut)- date as string
- text(ValueOut) - reducer writes max and min temperature with time as string
- public void reduce(**Text key, Iterable<Text> values, Context context**)
- Text key is value of mapper output. i.e:- City & date information
- Iterable<Text> values - values stores multiple temperature values for a given city and date.
- context object is where reducer write it's processed outcome and finally written in file.

MultipleOutputs :- In general, reducer generates output file(i.e: part_r_0000), however in this use case we want to generate multiple output files. In order to deal with such scenario we need to use MultipleOutputs of "org.apache.hadoop.mapreduce.lib.output.MultipleOutputs" which provides a way to write multiple file depending on reducer outcome. See below reducer class for more details. For each reducer task multipleoutput object is created and key/result is written to appropriate file.

Lets **create a Map/Reduce project in eclipse** and create a class file name it as CalculateMaxAndMinTemperatureWithTime. For simplicity, here we have written mapper and reducer class as inner static class. Copy following code lines and paste in newly created class file.

```
/**
 * Question:- To find Max and Min temperature from
 * record set stored in
 * text file. Schema of record set :-
 * tab separated (\t) CA_25-Jan-2014
 * 00:12:345 15.7 01:19:345 23.1
 * 02:34:542 12.3 03:12:187 16 04:00:093
 * -14 05:12:345 35.7 06:19:345 23.1
```

```
07:34:542 12.3 08:12:187 16
* 09:00:093 -7 10:12:345 15.7 11:19:345
23.1 12:34:542 -22.3 13:12:187
* 16 14:00:093 -7 15:12:345 15.7
16:19:345 23.1 19:34:542 12.3
* 20:12:187 16 22:00:093 -7
* Expected output:- Creates files for each city
and store maximum & minimum
* temperature for each day along
with time.
```

```
*/
```

```
import java.io.IOException ;
import java.util.StringTokenizer ;
import org.apache.hadoop.io.Text ;
import org.apache.hadoop.mapreduce.Mapper ;
import
org.apache.hadoop.mapreduce.Reducer ;
import
org.apache.hadoop.mapreduce.lib.output.MultipleOutput
ts ;
import org.apache.hadoop.conf.Configuration ;
import org.apache.hadoop.fs.Path ;
import org.apache.hadoop.mapreduce.Job ;
import
org.apache.hadoop.mapreduce.lib.input.FileInputForma
t ;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFor
mat ;
import
org.apache.hadoop.mapreduce.lib.output.TextOutputForm
at ;
/**
 * @author devinline
 */
public class CalculateMaxAndMinTemperatureWithTime {
public static String calOutputName =
```

```

"California";
public static String nyOutputName = "Newyork";
public static String njOutputName =
"Newjersy"; public static String ausOutputName
= "Austin"; public static String bosOutputName
= "Boston";
public static String balOutputName
= "Baltimore";
public static class WhetherForcastMapper
extends Mapper<Object, Text, Text,
Text> { public void map(Object
keyOffset, Text dayReport, Context
con)
throws IOException, InterruptedException {
StringTokenizer strTokens = new
StringTokenizer(
dayReport.toString(), "\\t");
int counter = 0;
Float currnetTemp = null;
Float minTemp =
Float.MAX_VALUE; Float maxTemp
= Float.MIN_VALUE; String date =
null;
String currentTime = null;
String minTempANDTime = null;
String maxTempANDTime = null;
while (strTokens.hasMoreElements()) {
if (counter == 0) {
date = strTokens.nextToken();
} else {
if (counter % 2 == 1) {
currentTime =

```

```

strTokens.nextToken(); } else {
currnetTemp =
Float.parseFloat(strTokens.nextToken());
if (minTemp > currnetTemp) {
minTemp = currnetTemp;
minTempANDTime = minTemp + "AND"
+ currentTime;
}
if (maxTemp < currnetTemp) {
maxTemp = currnetTemp;
maxTempANDTime = maxTemp + "AND"
+ currentTime;
}
}
}
counter++;
}
// Write to context - MinTemp, MaxTemp
and corresponding time
Text temp = new Text();
temp.set(maxTempANDTime);
Text dateText = new Text();
dateText.set(date);
try {
con.write(dateText, temp);
} catch (Exception e) {
e.printStackTrace();
}
temp.set(minTempANDTime);
dateText.set(date);
con.write(dateText, temp);
}
}

```

```

public static class WhetherForecastReducer
extends Reducer<Text, Text, Text,
Text> { MultipleOutputs<Text, Text>
mos; public void setup(Context context)
{ mos = new MultipleOutputs<Text,
Text>(context);
}
public void reduce(Text key,
Iterable<Text> values, Context
context)
throws IOException, InterruptedException {
int counter = 0;
String reducerInputStr[] = null;
String f1Time = "";
String f2Time = "";
String f1 = "", f2 = "";
Text result = new Text();
for (Text value : values) {
if (counter == 0) {
reducerInputStr =
value.toString().split("AND");
f1 = reducerInputStr[0];
f1Time = reducerInputStr[1];
}
else {
reducerInputStr =
value.toString().split("AND");
f2 = reducerInputStr[0];
f2Time = reducerInputStr[1];
}
counter = counter + 1;
}
}

```

```

if (Float.parseFloat(f1) >
Float.parseFloat(f2)) {
result = new Text("Time: " + f2Time +
" MinTemp: " + f2 + "\t"
+ "Time: " + f1Time + " MaxTemp: " +
f1); } else {
result = new Text("Time: " + f1Time +
" MinTemp: " + f1 + "\t"
+ "Time: " + f2Time + " MaxTemp: " +
f2); }
String fileName = "";
if (key.toString().substring(0,
2).equals("CA")) {
fileName =
CalculateMaxAndMinTemperatureTime.calOutputName
; } else if (key.toString().substring(0,
2).equals("NY")) {
fileName =
CalculateMaxAndMinTemperatureTime.nyOutputName
; } else if (key.toString().substring(0,
2).equals("NJ")) {
fileName =
CalculateMaxAndMinTemperatureTime.njOutputName
; } else if (key.toString().substring(0,
3).equals("AUS")) {
fileName =
CalculateMaxAndMinTemperatureTime.ausOutputName
; } else if (key.toString().substring(0,
3).equals("BOS")) {
fileName =
CalculateMaxAndMinTemperatureTime.bosOutputName
; } else if (key.toString().substring(0,

```

```

3).equals("BAL")) {
    fileName =
    CalculateMaxAndMinTemperatureTime.balOutputName
    ; }
    String strArr[] =
    key.toString().split("_");
    key.set(strArr[1]); //Key is date value
    mos.write(fileName, key, result);
}
@Override
public void cleanup(Context context)
throws IOException,
InterruptedException {
mos.close();
}
}

public static void main(String[] args)
throws IOException,
ClassNotFoundException,
InterruptedException {
    Configuration conf = new
    Configuration(); Job job =
    Job.getInstance(conf, "Wheather Statistics of
    USA");
    job.setJarByClass(CalculateMaxAndMinTemperature
    WithTime.class);
    job.setMapperClass(WhetherForecastMapper.class);
    job.setReducerClass(WhetherForecastReducer.class
    ) ;
    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(Text.class)
    ; job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

```

```
MultipleOutputs.addNamedOutput(job,
    calOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    nyOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    njOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    bosOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    ausOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    balOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
// FileInputFormat.addInputPath(job,
// new Path(args[0]));
// FileOutputFormat.setOutputPath(job,
// new Path(args[1]));
Path pathInput = new Path(
    "hdfs://192.168.213.133:54310/weatherInputData
/ input_temp.txt");
Path pathOutputDir = new Path(
    "hdfs://192.168.213.133:54310/user/hduser1
```



```

/ testfs/output_mapred3");
FileInputFormat.addInputPath(job,
pathInput);
FileOutputFormat.setOutputPath(job,
pathOutputDir);
try {
System.exit(job.waitForCompletion(true) ? 0
: 1);
} catch (Exception e) {
// TODO Auto-generated catch block
e.printStackTrace();
}
}
}
}
}

```

Explanation:-

In **map method**, we are parsing each input line and maintains a counter for extracting date and each temperature & time information. For a given input line, first extract date(counter ==0) and followed by alternatively extract time(counter%2==1) since time is on odd number position like (1,3,5....) and get temperature otherwise. Compare for max & min temperature and store it accordingly. Once while loop terminates for a given input line, write maxTempTime and minTempTime with date.

In **reduce method**, for each reducer task, setup method is executed and create MultipleOutput object. For a given key, we have two entry (maxtempANDTime and mintempANDTime). Iterate values list, split value and get temperature & time value. Compare temperature value and create actual value sting which reducer write in appropriate file.

In **main method**, a instance of Job is created with Configuration object. Job is configured with mapper, reducer class and along with input and output format. MultipleOutputs information added to Job to indicate file name to be used with input format. For this sample program, we are using input file("/weatherInputData/input_temp.txt") placed on HDFS and output directory (/user/hduser1/testfs/output_mapred5) will be also created on HDFS. Refer below command to copy downloaded input file from local file system to HDFS and give write permission to client who is executing this program unit so that output directory can be created.

Copy a input file form local file system to HDFS

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin\$

```
./hadoop fs -put /home/zytham/input_temp.txt  
/  
weatherInputData/
```

Give write permission to all user for creating output directory

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin\$./hadoop fs
-chmod -R 777 /user/hduser1/testfs/ Before
executing above program unit make sure hadoop services are running(to
start all service execute ./start-all.sh from <hadoop_home>/sbin). Now
execute above sample program. Run -> Run as hadoop. Wait for a moment
and check whether output directory is in place on HDFS. Execute following
command to verify the same.

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin\$./hadoop fs
-ls /user/hduser1/testfs/output_mapred3
Found 8 items
-rw-r--r-- 3 zytham supergroup 438
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **Austin-r-00000**
-rw-r--r-- 3 zytham supergroup 219
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **Baltimore-r-00000**
-rw-r--r-- 3 zytham supergroup 219
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **Boston-r-00000**
-rw-r--r-- 3 zytham supergroup 511
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **California-r-00000**
-rw-r--r-- 3 zytham supergroup 146
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **Newjersy-r-00000**
-rw-r--r-- 3 zytham supergroup 219
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ **Newyork-r-00000**
-rw-r--r-- 3 zytham supergroup 0
2015-12-11 19:21

```
/user/hduser1/testfs/output_mapred3/ _SUCCESS
-rw-r--r-- 3 zytham supergroup 0
2015-12-11 19:21
/user/hduser1/testfs/output_mapred3/ part-r-00000
```

Open one of the file and verify expected output schema, execute following command for the same.

```
hduser1@ubuntu:/usr/local/hadoop2.6.1/bin$ ./hadoop fs
-cat /user/hduser1/testfs/output_mapred3/
Austin-r-00000
```

```
25-Jan-2014 Time: 12:34:542 MinTemp: -22.3 Time:
05:12:345 MaxTemp: 35.7
```

```
26-Jan-2014 Time: 22:00:093 MinTemp: -27.0
Time: 05:12:345 MaxTemp: 55.7
```

```
27-Jan-2014 Time: 02:34:542 MinTemp: -22.3
Time: 05:12:345 MaxTemp: 55.7
```

```
29-Jan-2014 Time: 14:00:093 MinTemp: -17.0
Time: 02:34:542 MaxTemp: 62.9
```

```
30-Jan-2014 Time: 22:00:093 MinTemp: -27.0
Time: 05:12:345 MaxTemp: 49.2
```

```
31-Jan-2014 Time: 14:00:093 MinTemp: -17.0
Time: 03:12:187 MaxTemp: 56.0
```

Note:-

- In order to reference input file from local file system instead of HDFS, uncomment below lines in main method and comment below added addInputPath and setOutputPath lines. Here Path(args[0]) and Path(args[1]) **read input and output location path from program arguments**. OR create path object with sting input of input file and output location.

```
// FileInputFormat.addInputPath(job, new Path(args[0]));
// FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

Execute WeatherReportPOC.jar on single node

cluster We can create jar file out of this project and run on single node cluster too. **Download WeatherReportPOC jar** and place at some convenient location. Start hadoop services(./start-all.sh from <hadoop_home>/sbin). I have placed jar at

"/home/zytham/Downloads/WeatherReportPOC.jar". Execute following command to submit job with input file HDFS location is "/wheatherInputData/input_temp.txt" and output directory location is "/user/hduser1/testfs/output_mapred7"

hduser1@ubuntu:/usr/local/hadoop2.6.1/bin\$

./hadoop jar /home/zytham/Downloads/WeatherReportPOC.jar

**CalculateMaxAndMinTemperatureWithTime /
wheatherInputData/input_temp.txt /user/hduser1/
testfs/output_mapred7**

```
15/12/11 22:16:12 INFO
Configuration.deprecation: session.id is
deprecated. Instead, use
dfs.metrics.session-id
15/12/11 22:16:12 INFO jvm.JvmMetrics: Initializing
JVM Metrics with processName=JobTracker, sessionId=
15/12/11 22:16:14 WARN
mapreduce.JobResourceUploader: Hadoop command-line
option parsing not performed. Implement the Tool
interface and execute your application with
ToolRunner to remedy this. ....
15/12/11 22:16:26 INFO
output.FileOutputCommitter: Saved output of task
'attempt_local1563851561_0001_r_000000_0' to
hdfs://
hostname:54310/user/hduser1/testfs/output_mapred7/
_temporary/0/task_local1563851561_0001_r_000000
15/12/11 22:16:26 INFO mapred.LocalJobRunner: reduce
> reduce
15/12/11 22:16:26 INFO mapred.Task: Task
'attempt_local1563851561_0001_r_000000_0'
done. 15/12/11 22:16:26 INFO
mapred.LocalJobRunner:
Finishing task:
attempt_local1563851561_0001_r_000000_0
15/12/11 22:16:26 INFO mapred.LocalJobRunner:
```

```
reduce task executor complete.  
15/12/11 22:16:26 INFO mapreduce.Job: map  
100% reduce 100%  
15/12/11 22:16:27 INFO mapreduce.Job: Job  
job_local1563851561_0001 completed successfully  
15/12/11 22:16:27 INFO mapreduce.Job: Counters:  
38 .....
```